

ARTIFICIAL INTELLIGENCE LAB REPORT

Group Name: Acode

Tejas Pakhale-202211061, Rajat Kumar Thakur-202211070, Tanay Patel-202211094, Abhi Tundiya-202211095

Abstract—This comprehensive report explores a variety of Artificial Intelligence (AI) techniques, beginning with state-space search methods to solve classical problems such as the Missionaries and Cannibals and the Rabbit Leap. Both problems are modeled as state-space graphs, where Breadth-First Search (BFS) and Depth-First Search (DFS) are employed to compare optimality, memory usage, and computational trade-offs. Extending the exploration of search techniques, the A* algorithm is used to align text documents for plagiarism detection and solve the 8-puzzle problem, focusing on time and memory requirements.

The report then delves into non-deterministic search methods, applying heuristic-driven algorithms like Hill-Climbing, Beam Search, and Variable-Neighborhood Descent to solve complex problems such as marble solitaire and random k-SAT generation. Simulated Annealing is applied to the Traveling Salesman Problem (TSP), showcasing its utility in probabilistic optimization and tackling large search spaces.

In the domain of probabilistic graphical models, the report implements Bayesian Networks to explore relationships between student course grades, employing naive Bayes classifiers for internship qualification prediction. The classifier's accuracy is tested using repeated experiments, comparing independent and dependent variable assumptions across training and testing datasets. Through these diverse AI techniques and experiments, the report highlights the strengths and applications of search algorithms, heuristic functions, and probabilistic models in solving complex real-world problems.

I. INTRODUCTION

This report delves into fundamental AI concepts through hands-on experiments that explore state-space search, heuristic-driven optimization, and probabilistic models. By implementing algorithms such as Breadth-First Search (BFS), Depth-First Search (DFS), A*, and Simulated Annealing, participants tackle classical problems like the Missionaries and Cannibals, the 8-puzzle, and the Traveling Salesman Problem. Additionally, the report covers probabilistic reasoning through Bayesian Networks, providing insights into AI's capacity for decision-making under uncertainty.

Week 1: Modeling and Solving Classical Search Problems Using State Space Search Techniques

To be able to model a given problem in terms of state space search problem and solve the same using BFS/ DFS.

Week 2: State Space Search

To design a graph search agent and understand the use of a hash table, queue in state space search.

Week 3: Heuristic And Non-Classical Search

To understand the use of Heuristic function for reducing the size of the search space. Explore non-classical search

algorithms for large problems.

Week 4: Non-deterministic Search and Simulated Annealing

Non-deterministic Search | Simulated Annealing For problems with large search spaces, randomized search becomes a meaningful option given partial/ full-information about the domain.

Week 5: Graphical Models and Bayesian Networks Inference and Classification in R

Understand the graphical models for inference under uncertainty, build Bayesian Network in R, Learn the structure and CPTs from Data, naive Bayes classification with dependency between features.

Challenge Problem : Generating melody using simulated annealing

Generate a melody in North Indian Classical Raag Bhairav using either simulated annealing or the genetic algorithm

II. WEEK 1: MODELING AND SOLVING CLASSICAL SEARCH PROBLEMS USING STATE SPACE SEARCH TECHNIQUES

A. Introduction

State space search techniques are particularly valuable for addressing problems, where the objective is to navigate through different configurations or states to arrive at a solution. This report examines two such problems—the Missionaries and Cannibals and the Rabbit Leap—which have been widely used as benchmarks in AI research.

The Missionaries and Cannibals involves a group of missionaries and cannibals who must be safely transported across a river using a boat, with the constraint that at no point should the cannibals outnumber the missionaries on either side.

The Rabbit Leap problem presents a different type of challenge, involving two groups of rabbits—one moving eastward and the other westward—who must cross each other along a sequence of stones placed in a straight line across a stream. The puzzle requires the rabbits to move forward or leap over one another while following the rules of the game. The challenge is to determine whether the rabbits can cross each other without violating the movement constraints or falling into the water.

Both problems serve as ideal case studies for the application of search techniques like Breadth-First Search (BFS) and Depth-First Search (DFS), allowing for a detailed analysis of search efficiency, optimality, and computational complexity in solving such puzzles.

B. UNDERSTANDING

1) **Search Algorithms:** Search algorithms provide systematic approaches to explore a state space, navigating through possible configurations to reach a solution. They are vital in artificial intelligence (AI) for addressing complex problems by transitioning from an initial state to a goal state based on predefined rules and constraints. Before diving into practical applications, it is important to understand the types of search algorithms, such as *Breadth-First Search (BFS)* and *Depth-First Search (DFS)*, and the distinction between *uninformed* and *informed* search strategies, including the role of heuristics in the latter.

2) **State Space:** A state space is the set of all possible configurations or arrangements that a problem can take. It forms the foundation for search techniques, as each state represents a unique point in the problem's progression.

Researchers must first define the state space carefully for problems like *Missionaries and Cannibals* or *Rabbit Leap*, outlining each possible combination of missionaries, cannibals, boats, or rabbits that could occur throughout the puzzle.

3) **Search Efficiency:** Search efficiency refers to how effectively a search algorithm can find the solution, typically measured in terms of time complexity (how long the algorithm takes) and space complexity (the amount of memory it uses).

Analyzing search efficiency involves comparing algorithms like BFS and DFS, which differ in how they explore the state space. Factors such as the depth of the solution and the branching factor of the problem significantly impact efficiency.

4) **Optimality:** Optimality ensures that a search algorithm finds not just any solution, but the best possible one. For instance, in puzzles like the *Rabbit Leap* problem, an optimal solution would be the one that requires the least moves or satisfies the goal with minimal cost.

The choice of the search algorithm is largely dependent on the constraint of the problem as well as the available system resources. Researchers focus on search algorithms like BFS, which guarantees an optimal solution in terms of the shortest path, compared to DFS, which may find a solution but not necessarily the optimal one while they may also opt for DFS when time is a constraint, especially if the solution obtained is close to the optimal one and the difference is not significant. This approach is practical when the speed of finding a solution outweighs the need for exact optimality.

C. Problem Statement

1) **Missionaries and Cannibals:** The missionaries and cannibals problem is usually stated as follows. Three missionaries and three cannibals are on one side of a river, along with a boat that can hold one or two people. Find a way to get everyone to the other side without ever leaving a group of missionaries in one place outnumbered by the cannibals in that place.

2) **The Rabbit Leap:** In the rabbit leap problem, three east-bound rabbits stand in a line blocked by three west-bound rabbits. They are crossing a stream with stones placed in the east west direction in a line. There is one empty stone between them. The rabbits can only move forward one or two steps. They can jump over one rabbit if the need arises, but not more than that.

3) Problem I:

- Model the problem as a state space search problem. How large is the search space?
- Solve the problem using BFS. The optimal solution is the one with the fewest number of steps. Is the solution that you have acquired an optimal one? The program should print out the solution by listing a sequence of steps needed to reach the goal state from the initial state.
- Solve the problem using DFS. The program should print out the solution by listing a sequence of steps needed to reach the goal state from the initial state.
- Compare solutions found from BFS and DFS. Comment on solutions. Also compare the time and space complexities of both.

D. Missionaries and Cannibals Problem

1) **Problem Formulation:** We represent the state as a tuple (M, C, B) where:

- M : Number of missionaries on the left side
- C : Number of cannibals on the left side
- B : Position of the Boat (1: Left bank, 2: Right Bank)

Initial State: (3, 3, 1)

Goal State: (0, 0, 0)

For missionaries, we have 4 possibilities

(0, 1, 2, or 3 missionaries)

For cannibals, we also have 4 possibilities

(0, 1, 2, or 3 cannibals)

For the boat, we have 2 possibilities

(0 or 1)

$$\Rightarrow 4 \times 4 \times 2 = 32 \text{ possible states} \quad (1)$$

2) BFS Solution:

Algorithm 1 Breadth first Space for Missionaries and Cannibals Problem

```

1: procedure BFS(start_state, goal_state)
2:   queue  $\leftarrow$  deque([(start_state, [])])
3:   visited  $\leftarrow$   $\emptyset$ 
4:   while queue is not empty do
5:     (state, path)  $\leftarrow$  queue.popleft()
6:     if state  $\in$  visited then
7:       continue
8:     end if
9:     nodes_visited  $\leftarrow$  nodes_visited + 1
10:    path  $\leftarrow$  path + [state]
11:    if state = goal_state then
12:      print nodes_visited
13:      return path
14:    end if
15:    for successor  $\in$  get_successors(state) do
16:      queue.append((successor, path))
17:    end for
18:  end while
19:  return None
20: end procedure

```

3) *DFS Solution:***Algorithm 2** Depth-First Search for Missionaries and Cannibals Problem

```

1: procedure DFS(start_state, goal_state)
2:   stack  $\leftarrow$  deque([(start_state, [])])
3:   visited  $\leftarrow$   $\emptyset$ 
4:   while stack is not empty do
5:     (state, path)  $\leftarrow$  stack.pop()
6:     if state  $\in$  visited then
7:       continue
8:     end if
9:     visited.add(state)
10:    path  $\leftarrow$  path + [state]
11:    if state = goal_state then
12:      return path
13:    end if
14:    for successor  $\in$  get_successors(state) do
15:      stack.append((successor, path))
16:    end for
17:  end while
18:  return None
19: end procedure

```

Complexity Type	BFS	DFS
Time Complexity	$O(b^d)$	$O(b^d)$
Space Complexity	$O(b^d)$	$O(bd)$
Total visited	15	14

Table I: Complexity Comparison of BFS and DFS

4) *Complexity Comparison:* where b is branching factor and d is depth of shallowest goal state.

E. Analysis

The time complexity of the BFS algorithm is $O(b^d)$, where b is the branching factor (the average number of successors per state), and d is the depth of the shallowest goal state. In this case, the `get_successors` function generates all possible successors for a given state, contributing to the branching factor. The BFS function's while loop iterates through the queue of states until the goal is reached or all nodes are visited, resulting in a time complexity of $O(b^d)$. Similarly, the time complexity of the DFS algorithm is also $O(b^d)$, where b is the branching factor (number of possible moves from each state), and d is the depth of the search tree. In the Rabbit Leap Problem, the branching factor is 4 (1 step or 2 steps in either direction), while the depth varies depending on the specific instance of the problem.

Depth-First Search (DFS) is more space-efficient than Breadth-First Search (BFS) because it explores one path to its maximum depth before backtracking, only needing to store the nodes along the current path. This leads to a space complexity of $O(bd)$, where b is the branching factor and d is the maximum depth of the tree. In contrast, BFS explores all nodes at each depth level before proceeding, requiring it to store all nodes at a given depth, resulting in a space complexity of $O(b^d)$, where d is the depth of the shallowest solution.

In large or wide search spaces, BFS can quickly consume exponentially more memory, making DFS far more space-efficient when memory resources are limited. For the Rabbit Leap problem, the state space is quite structured, meaning that DFS might find the goal faster than BFS due to its depth-first nature. Since BFS explores all nodes at the same depth level before moving deeper, it tends to visit more nodes, especially if the solution is located at a deeper level in the tree. DFS, on the other hand, can reach the solution sooner if it explores the right branches first.

F. Conclusion

In conclusion, this laboratory highlights the comparative analysis of Breadth-First Search (BFS) and Depth-First Search (DFS) algorithms in solving the Rabbit Leap Problem. While both algorithms exhibit a time complexity of $O(b^d)$, DFS demonstrates greater space efficiency due to its depth-focused approach. The choice between BFS and DFS ultimately depends on the specific characteristics of the problem, including the branching factor and depth of the search tree. In the given set of problem the DFS was superior compared to BFS in all aspect.

III. WEEK 2: STATE SPACE SEARCH

A. Introduction

This lab report consists of graph search algorithms to solve two distinct problems: plagiarism detection in text documents and the 8-puzzle. The goal is to design a graph search agent using techniques like A* and Iterative Deepening Search. By exploring these algorithms, we analyze their effectiveness in navigating problem spaces and their implications for memory and time efficiency.

The 8-puzzle is a classic AI problem that consists of a 3x3 grid with eight numbered tiles and one empty space. The objective is to rearrange the tiles from an initial configuration to a goal state by sliding them one at a time into the empty space. To solve this, we apply state-space search algorithms, such as A* and Iterative Deepening Search, to find the optimal sequence of moves that achieves the goal in the fewest steps.

The plagiarism detection task involves comparing two documents to identify similar or identical sequences of text, which may indicate copying. We implement the A* search algorithm to align sentences between the documents by minimizing the edit distance, such as Levenshtein distance. This approach helps to detect potential instances of plagiarism by finding closely matching text segments.

*B. Puzzle-8*1) *Problem Statement:*

- 1) Write a pseudocode for a graph search agent. Represent the agent in the form of a flow chart. Clearly mention all the implementation details with reasons.
- 2) Write a collection of functions imitating the environment for Puzzle-8.
- 3) Describe what is Iterative Deepening Search.

- 4) Considering the cost associated with every move to be the same (uniform cost), write a function which can backtrack and produce the path taken to reach the goal state from the source/ initial state.
 - 5) Generate Puzzle-8 instances with the goal state at depth "d".
 - 6) Prepare a table indicating the memory and time requirements to solve Puzzle-8 instances (depth "d") using your graph search agent.
- 2) **Solution:**
- 1) The pseudocode for the following problem can be accessed through this [link](#) and the flowchart of the following is:

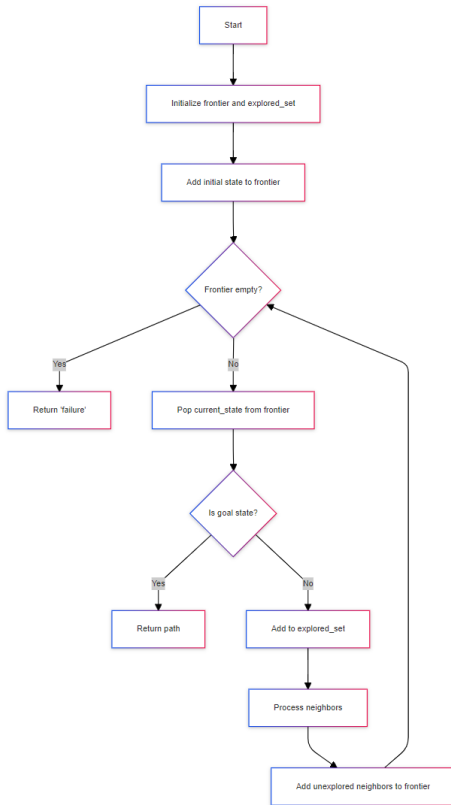


Figure 1: Flowchart for the Graph Search Agent

- 2) Collection of functions are:

- **Initial State:** Any state can be the initial state where the configuration of blocks is stored in a 3×3 matrix.
- **Actions:** Instead of moving the numbered blocks, we can consider the action as the blank space moving up, down, right, and left.
- **Goal State:** Although any state could be the goal, the correct configuration of blocks (the expected final state) is considered the goal state.
- **Goal Test:** Checking whether the new state is the same as the goal state or not.
- **Transitions:** Switching the positions of moved and empty blocks, resulting in a new state.
- **Path Cost:** Every step will have a cost, which will be added to the total cost.

- 3) Iterative deepening search involves applying depth-limited search multiple times, incrementing the depth limit with each iteration to efficiently locate a solution. In every iteration, there are three potential outcomes: finding a solution node, encountering failure if no solution exists within the current depth, or hitting a cutoff point that suggests a possible solution at a deeper level. This method does not keep track of previously reached states, which means the same state might be revisited multiple times along different paths. The optimal depth limit is determined by progressively testing limits starting from 0, then 1, and so on, until a solution is found. This ensures a comprehensive exploration of the state space. The time complexity of iterative deepening search depends on whether a solution is present. If a solution is found, the time complexity is $O(b^d)$, where b is the branching factor and d is the solution depth. If no solution exists, the complexity becomes $O(b^m)$, where m is the maximum explored depth.
- 4) The pseudo-code for the following problem can be accessed through [link](#)
- 5) The pseudo-code for the following problem can be accessed through [link](#)
- 6) The table is as follows:

Depth	Nodes Explored	Time Elapsed (s)	Memory Used (bytes)
1	12.00	10.0000	10
2	12.60	10.0000	1819
3	12.80	10.0000	10
4	13.00	10.0003	10
5	12.00	10.0000	10
6	13.20	10.0000	10
7	13.60	10.0002	10
8	17.80	10.0000	12458
9	14.40	10.0000	10
10	13.80	10.0002	10
11	16.60	10.0000	10
12	16.40	10.0002	10
13	17.00	10.0000	10
14	129.00	10.0002	12288
15	16.20	10.0002	10
16	17.40	10.0000	10
17	14.00	10.0002	10
18	12.40	10.0002	10
19	19.00	10.0000	10
20	131.40	10.0002	10

Figure 2: Table indicating memory and time requirements

C. Plagiarism Detection System

1) **Problem Statement:** Develop a plagiarism detection system that identifies similar or identical sequences of text between two documents by aligning their sentences or paragraphs. The system will leverage the A* search algorithm to efficiently find optimal text alignments based on edit distance, indicating potential instances of plagiarism.

2) **Levenshtein Distance:** Levenshtein distance is a similarity measure between two strings, representing the minimum number of edit operations (insertions, deletions, or substitutions) required to transform one string into another. For example, changing the word "test" into "tent" requires one substitution, so the Levenshtein distance is 1.

In the context of this lab, the Levenshtein distance is used to compute the cost of aligning sentences from two documents, which is key to detecting plagiarism. By aligning text based on minimal edit distance, the system identifies similar or identical sequences of text.

Algorithm and Use in Plagiarism Detection

Given two text documents, the system aligns their sentences using the A* search algorithm. The alignment minimizes the Levenshtein distance between corresponding sentences. The Levenshtein distance provides the following advantages in this context:

- **Cost Function (g(n)):** The total edit distance between aligned sentences is used to compute the cost. Smaller distances indicate higher similarity and possible plagiarism.
- **Heuristic Function (h(n)):** The remaining sentences' minimum edit distances are estimated to guide the A* search algorithm to an optimal alignment efficiently.

In this system, a lower Levenshtein distance between two sentences indicates a higher degree of similarity, which may be indicative of plagiarism. Conversely, a high edit distance suggests that the sentences differ significantly, reducing the likelihood of plagiarism.

Example of Levenshtein Distance

Consider two words: "GUMBO" and "GAMBOL". The Levenshtein distance is 2 because transforming "GUMBO" into "GAMBOL" requires one substitution ("U" to "A") and one insertion ("L"). This concept scales to entire sentences in plagiarism detection, where the number of edits required to transform one sentence into another helps identify potentially copied content.

		G	U	M	B	O
	0	1	2	3	4	5
G	1	0	1	2	3	4
A	2	1	1	2	3	4
M	3	2	2	1	2	3
B	4	3	3	2	1	2
O	5	4	4	3	2	1
L	6	5	5	4	3	2

Figure 3: Levenshtein Distance Matrix for "GUMBO" and "GAMBOL"

3) **Solution:** The code for this problem statement can be accessed through this [link](#).

Here's an overview of the steps taken to arrive at the solution:

- 1) **Text Preprocessing:** A 'preprocess_text' function was created that:
 - Splits the input text into sentences using regular expressions.
 - Converts each sentence to lowercase.
 - Removes leading and trailing whitespace from each sentence.
- 2) **Similarity Calculation:** To calculate similarity between sentences, the Levenshtein distance was utilized:

- A 'levenshtein_distance' function was implemented to compute the edit distance between two strings.
- A 'calculate_similarity' function was developed that normalizes the Levenshtein distance to a percentage similarity.

3) **Plagiarism Detection:** The 'detect_plagiarism' function was created to tie everything together:

- Finds the best pairs of sentences between the documents using A* algorithm to align sentences in both documents.
- Calculates an overall similarity percentage for the entire documents.

The Levenshtein distance between two strings s_1 and s_2 is given by the following recurrence relation:

$$lev_{s_1, s_2}(i, j) = \begin{cases} \max(i, j) \\ \min \begin{cases} lev_{s_1, s_2}(i-1, j) + 1 \\ lev_{s_1, s_2}(i, j-1) + 1 \\ lev_{s_1, s_2}(i-1, j-1) + 1_{(s_1[i] \neq s_2[j])} \end{cases} \end{cases}$$

where $1_{(s_1[i] \neq s_2[j])}$ is the indicator function equal to 0 when $s_1[i] = s_2[j]$ and equal to 1 otherwise.

The percentage similarity between two strings can be calculated using the following formula:

$$\text{similarity}(s_1, s_2) = \left(1 - \frac{\text{levenshtein_distance}(s_1, s_2)}{\max(|s_1|, |s_2|)}\right) \times 100$$

Where: - $\text{levenshtein_distance}(s_1, s_2)$ is the Levenshtein distance between the two strings s_1 and s_2 . - $|s_1|$ and $|s_2|$ are the lengths of the strings s_1 and s_2 respectively.

4) **A* Algorithm for document alignment:** **Input:** Two documents doc1 and doc2.

Output: Optimal alignment between sentences

1) Initialization:

- Create the initial state with $\text{pos1} = 0$, $\text{pos2} = 0$, $\text{cost} = 0$, and an empty path.
- Initialize a priority queue (min-heap) with the initial state.
- Initialize an empty set to keep track of visited positions.

2) Main Loop:

- While the priority queue is not empty:
 - Pop the state with the lowest cost from the queue.
 - If both $\text{pos1} = |\text{doc1}|$ and $\text{pos2} = |\text{doc2}|$, return the alignment path (goal reached).
 - If $(\text{pos1}, \text{pos2})$ has already been visited, continue to the next iteration.
 - Mark $(\text{pos1}, \text{pos2})$ as visited.

3) State Expansion:

- If both $\text{pos1} < |\text{doc1}|$ and $\text{pos2} < |\text{doc2}|$
 - Compute the cost using the Levenshtein distance between $\text{doc1}[\text{pos1}]$ and $\text{doc2}[\text{pos2}]$.
 - Create a new state advancing both pos1 and pos2 , update the total cost, and append the alignment to the path.

- Push the new state onto the priority queue with priority equal to the cost plus a heuristic.
- If $\text{pos1} < |\text{doc1}|$ (skip sentence in doc2):
 - Create a new state advancing pos1 only, with an additional cost of 1, and push it onto the queue.
- If $\text{pos2} < |\text{doc2}|$ (skip sentence in doc1):
 - Create a new state advancing pos2 only, with an additional cost of 1, and push it onto the queue.

4) Heuristic:

- Use a heuristic function to estimate the remaining alignment cost by taking the minimum of the remaining sentences between doc1 and doc2:

$$h(\text{state}) = \min(|\text{doc1}| - \text{pos1}, |\text{doc2}| - \text{pos2})$$

5) *Result of Test-Cases:* The result of test cases can be seen below:

```
Test Case: Identical Documents
Document 1: This is a test. It has multiple sentences. We want to detect plagiarism.
Document 2: This is a test. It has multiple sentences. We want to detect plagiarism.
Detected 3 plagiarism cases:
  Sentence 1 in doc1 matches sentence 1 in doc2 with 100.00% similarity
  Sentence 2 in doc1 matches sentence 2 in doc2 with 100.00% similarity
  Sentence 3 in doc1 matches sentence 3 in doc2 with 100.00% similarity
Overall similarity: 100.00%
Verdict: Surely copied.

Test Case: Slightly Modified Document
Document 1: This is a test. It has multiple sentences. We want to detect plagiarism.
Document 2: This is an exam. It contains several phrases. We aim to identify copying.
Detected 0 plagiarism cases:
Overall similarity: 46.23%
Verdict: Possibly copied.
```

```
Test Case: Completely Different Documents
Document 1: This is about cats. Cats are furry animals. They make good pets.
Document 2: Python is a programming language. It is widely used in data science.
Detected 0 plagiarism cases:
Overall similarity: 19.93%
Verdict: Not copied.

Test Case: Partial Overlap
Document 1: This is a test. It has multiple sentences. We want to detect plagiarism.
Document 2: This is different. We want to detect plagiarism. This is unique.
Detected 1 plagiarism cases:
  Sentence 3 in doc1 matches sentence 2 in doc2 with 100.00% similarity
Overall similarity: 65.24%
Verdict: Likely copied.
```

Figure 4: Test Cases Verdict

D. Conclusion

In conclusion, this lab report presents a comprehensive exploration of graph search techniques through the implementation of a graph search agent for the Puzzle-8 problem, along with a detailed examination of Iterative Deepening Search. By developing functions to simulate the Puzzle-8 environment and implementing a backtracking function for path retrieval, we demonstrated the effectiveness of uniform cost searches. The generation of Puzzle-8 instances at various depths facilitated an analysis of memory and time requirements, providing valuable insights into the performance of our search algorithms. Additionally, we extended our investigation to implement a

plagiarism detection system using the A* search algorithm applied to text alignment. This system successfully identifies similar or identical text sequences between documents, enhancing the capability to detect potential plagiarism through effective alignment of sentences or paragraphs.

IV. WEEK 3: HEURISTIC AND NON-CLASSICAL SEARCH

A. Introduction

Tackling complex problems like marble solitaire and **Boolean satisfiability (SAT)** problems is a significant challenge in computer science. Marble solitaire involves removing marbles from a board until only one remains, while **k-SAT problems** focus on determining satisfying assignments for Boolean variables across multiple clauses. Both of these problems fall under the category of NP-hard problems, meaning they are notoriously difficult to solve efficiently for larger instances.

This experiment aims to explore non-deterministic search methods, augmented by heuristic functions, to find optimal or near-optimal solutions to these problems. The study of heuristic and non-classical search techniques is not only intellectually stimulating but also highly relevant in real-world applications. For example, the principles used to solve marble solitaire can inform various game AI algorithms, while k-SAT problems are foundational in optimizing network designs, scheduling tasks, and other computational challenges.

B. Non-deterministic Search

Non-deterministic search algorithms offer a flexible framework for exploring multiple potential solutions without following a predetermined path. Unlike deterministic algorithms, which yield consistent results for the same input, non-deterministic approaches introduce randomness at various stages of the search process. This is especially useful for navigating large search spaces where exhaustive searches become computationally expensive or impractical.

In the context of marble solitaire, non-deterministic search allows us to explore various sequences of moves that could lead to the desired configuration. Similarly, when tackling k-SAT problems, search algorithms can investigate different assignments of Boolean variables to find a combination that satisfies all clauses. The element of randomness in the search helps to prevent the algorithm from becoming trapped in local optima, enabling it to uncover more promising areas within the search space.

C. Marble Solitaire Problem

Marble solitaire is a classic puzzle where the objective is to remove all but one marble, ideally leaving the last marble in the center of the board. The game is played by jumping one marble over another into an empty hole, effectively removing the jumped marble from the board. This process continues until only one marble remains.

To solve marble solitaire, we employ the following strategies:

- 1) Implement a priority queue-based search considering path cost.
- 2) Propose two distinct heuristic functions, with justifications.
- 3) Implement the Best First Search algorithm.
- 4) Implement the A* algorithm.
- 5) Compare the results of these various search algorithms.

1) *Algorithms for Marble Solitaire: 1. Priority Queue-Based Search*

Input: Initial board configuration and a set of valid moves.

Output: Sequence of moves leading to the goal state.

Algorithm:

Algorithm 3 Priority Queue Search

```

1: Initialize a priority queue
2: Add the initial state to the queue with priority 0
3: while the queue is not empty do
4:   Pop the state with the highest priority
5:   if the goal state is reached then
6:     Return the path to the goal
7:   end if
8:   for each valid move do
9:     Generate a new state
10:    Calculate the path cost
11:    Add the new state to the queue with the calculated
    priority
12:   end for
13: end while
14: Return failure

```

Pros: This method guarantees that the first solution found is optimal, as it explores the most promising nodes first.

Cons: The priority queue can grow large, leading to high memory usage and longer search times, especially as the state space expands.

2. *Heuristic Functions:* We propose two heuristic functions to enhance our search:

- **Manhattan Distance Heuristic:** This measures how far each marble is from the center, providing a straightforward metric for guiding the search. By summing the distances of all marbles from their target positions, we gain insight into the potential moves needed to reach the goal.
- **Exponential Distance Heuristic:** This is similar, but with a larger bias towards the center. If H and V are the horizontal and vertical distances from the center respectively, then the heuristic's value is $2^{\max(H, V)}$.

Pros of Heuristic Functions: Heuristics can significantly reduce the search space, leading to faster solutions.

Cons: Poorly designed heuristics may misguide the search, resulting in longer solution times.

3. *Best First Search:* **Input:** Initial board configuration and a set of valid moves.

Output: Optimal sequence of moves leading to the goal state.

Algorithm:

Algorithm 4 Best First Search

```

1: Initialize an empty priority queue
2: Add the initial state with cost using the heuristic
3: while the queue is not empty do
4:   Dequeue the state with the lowest cost
5:   if the goal state is reached then
6:     Return the path to the goal
7:   end if
8:   for each valid move do
9:     Generate a new state
10:    Calculate the cost using the heuristic
11:    Enqueue the new state with the updated cost
12:   end for
13: end while
14: Return failure

```

Pros: This approach efficiently finds solutions by employing an informed search strategy that takes the estimated cost into account.

Cons: It may be slower than other methods if the heuristic is poorly tuned, as it might explore suboptimal paths.

4. *A* Search Algorithm:* **Input:** Initial board configuration and a set of valid moves.

Output: Optimal sequence of moves leading to the goal state.

Algorithm:

Algorithm 5 A* Search Algorithm

```

1: Initialize an empty priority queue
2: Add the initial state with cost 0
3: while the queue is not empty do
4:   Dequeue the state with the lowest  $f(n) = g(n) + h(n)$ 
5:   if the goal state is reached then
6:     Return the path to the goal
7:   end if
8:   for each valid move do
9:     Generate a new state
10:    Calculate  $g(n)$  and  $h(n)$ 
11:    Enqueue the new state with  $f(n)$ 
12:   end for
13: end while
14: Return failure

```

Pros: A* is highly efficient because it combines the cost to reach a node with the estimated cost to the goal, leading to a more directed search.

Cons: It requires more memory compared to simpler algorithms, as it keeps track of all nodes in the search space.

Results and Conclusion: After implementing these algorithms, we compared their performance based on several criteria: the number of moves required, time complexity, and overall efficiency in reaching the goal state. It's expected that the A* search algorithm will outperform the others due to its informed nature, while the best-first search offers a good balance between performance and simplicity. The priority queue search, while optimal, may lag in efficiency due to its higher memory requirements.

D. *k*-SAT Problem

The *k*-SAT problem is a well-known NP-complete problem that involves determining if a Boolean formula in conjunctive normal form (CNF) is satisfiable. In this experiment, we focus on the random generation of 3-SAT problems and solving them using various search algorithms.

A. Generating *k*-SAT Problems: To generate random *k*-SAT problems, we follow these steps:

- 1) Choose a number of variables n .
- 2) Choose a number of clauses m .
- 3) Randomly generate m clauses where each clause contains 3 literals, selecting from the set of variables and their negations.

Input: Number of variables n , number of clauses m .

Output: A random *k*-SAT problem in CNF.

Algorithm:

Algorithm 6 Generate Random 3-SAT Problem

- 1: Initialize an empty list of clauses
 - 2: **for** $i = 1$ to m **do**
 - 3: Create a clause with 3 random literals
 - 4: Add the clause to the list
 - 5: **end for**
 - 6: Return the list of clauses
-

B. Solving *k*-SAT Problems: When it comes to solving *k*-SAT problems, we employ three algorithms:

- 1) Hill-Climbing Algorithm
- 2) Beam Search Algorithm
- 3) Variable-Neighborhood-Descent Algorithm

1) **1. Hill-Climbing Algorithm:** **Input:** A random 3-SAT problem.

Output: A satisfying assignment or failure.

Algorithm:

Algorithm 7 Hill-Climbing Algorithm

- 1: Initialize a random assignment of variables
 - 2: **while** not satisfied **do**
 - 3: **if** current assignment satisfies the formula **then**
 - 4: Return current assignment
 - 5: **end if**
 - 6: Select a variable to flip
 - 7: Generate new assignment
 - 8: **if** new assignment is better **then**
 - 9: Update current assignment
 - 10: **end if**
 - 11: **end while**
 - 12: Return failure
-

Pros: The hill-climbing algorithm is straightforward and easy to grasp, making it a solid choice for those new to solving SAT problems.

Cons: However, it often gets stuck in local optima, which means it may miss the global solution.

2. Beam Search Algorithm: **Input:** A random 3-SAT problem.

Output: A satisfying assignment or failure.

Algorithm:

Algorithm 8 Beam Search Algorithm

- 1: Initialize beam with random assignments
 - 2: **while** not satisfied **do**
 - 3: Expand each assignment in the beam
 - 4: Evaluate all new assignments
 - 5: Keep the best w assignments in the beam
 - 6: **end while**
 - 7: Return best assignment or failure
-

Pros: Beam search effectively narrows down the search space by maintaining only the best candidates, which can enhance performance.

Cons: The effectiveness of the solution is highly dependent on the beam width; if the beam is too narrow, it may overlook optimal solutions.

3. Variable-Neighborhood-Descent Algorithm: **Input:** A random 3-SAT problem.

Output: A satisfying assignment or failure.

Algorithm:

Algorithm 9 Variable-Neighborhood-Descent Algorithm

- 1: Initialize random assignment of variables
 - 2: **while** not satisfied **do**
 - 3: **for** each neighborhood function **do**
 - 4: Generate new assignment
 - 5: **if** new assignment satisfies the formula **then**
 - 6: Return new assignment
 - 7: **end if**
 - 8: **end for**
 - 9: Update current assignment based on the best found
 - 10: **end while**
 - 11: Return failure
-

Pros: This algorithm expands the search space systematically by exploring various neighborhoods, which helps escape local optima.

Cons: It may require more iterations than simpler methods, potentially affecting overall efficiency.

Comparative Analysis: After solving the uniform random 3-SAT problems using the three different algorithms, we compare their performance based on:

- The number of iterations to reach a solution.
- The quality of solutions found.
- Time complexity.

It's anticipated that the Beam Search and Variable-Neighborhood-Descent algorithms will outperform the Hill-Climbing algorithm in terms of solution quality, especially as the complexity of SAT instances increases. This comparative analysis highlights the practical effectiveness of various problem-solving approaches.

E. Conclusion

In summary, both marble solitaire and k-SAT problems pose substantial challenges that can be effectively addressed using non-deterministic search techniques and heuristics. Algorithms like A* and Beam Search significantly enhance search efficiency, enabling exploration of larger search spaces. Our findings illustrate that informed search strategies can greatly improve performance in solving complex problems, paving the way for further research into optimization and algorithm development. This exploration not only enriches our theoretical understanding of these challenges but also promotes practical applications across fields like artificial intelligence, operations research, and game development.

V. WEEK 4: NON-DETERMINISTIC SEARCH AND SIMULATED ANNEALING

A. Introduction

The Travelling Salesman Problem (TSP) is a key challenge in combinatorial optimization. It asks: "Given a set of cities and distances, what is the shortest route that visits each city once and returns to the start?" The TSP is crucial in computer science and operations research because it is NP-hard, meaning solving it efficiently for large inputs is extremely difficult, and no known algorithm can solve all cases quickly.

The Travelling Salesman Problem (TSP) can be visualized as an undirected weighted graph, where cities are the vertices and the paths between them are the edges, each with a distance or cost. The goal is to find the shortest Hamiltonian cycle, a route that visits every vertex exactly once and returns to the starting point, minimizing the total travel distance. TSP is usually modeled as a complete graph, where every pair of cities is connected by an edge. If no direct path exists between cities, an artificial edge with a large weight can be added to ensure the graph remains complete without affecting the optimal solution.

The Travelling Salesman Problem (TSP) has practical applications in fields like logistics, transportation, telecommunications, and circuit design. Despite its computational challenges, researchers continually develop new algorithms and heuristics to find efficient, near-optimal solutions for real-world problems, such as route planning, network optimization, and supply chain management.

B. NON-DETERMINISTIC SEARCH

Non-deterministic search is a problem-solving approach that explores multiple potential solutions simultaneously, allowing algorithms to make random moves at each step. This flexibility enables the discovery of various solutions, particularly in complex scenarios like the Travelling Salesman Problem, where exhaustive searches for optimal solutions can be time-consuming. Unlike deterministic algorithms, which yield consistent outputs, non-deterministic algorithms can exhibit different behaviors across executions, investigating multiple routes to uncover a range of outcomes. This variability enhances their effectiveness in optimization tasks, particularly when exact

solutions are computationally expensive or unrealistic, facilitating innovative solutions that might not be found through traditional methods.

C. TRAVELLING SALESMAN PROBLEM

The Travelling Salesman Problem (TSP) is a well-known question in combinatorial optimization that asks: "Given a list of cities and the distances between each pair, what is the shortest route that visits each city exactly once and returns to the starting point?" As an NP-hard problem, the TSP is significant in both theoretical computer science and operations research, as it exemplifies the challenges of optimizing routes and logistics in various practical applications.

Mathematical Formulation:

$$\text{Minimize } Z = \sum_{i=1}^n \sum_{j=1}^n d_{ij} x_{ij}$$

Constraints:

$$\sum_{j=1}^n x_{ij} = 1 \quad \forall i \quad (1 \leq i \leq n)$$

$$\sum_{i=1}^n x_{ji} = 1 \quad \forall j \quad (1 \leq j \leq n)$$

- Z : Represents the total distance (or cost) of the tour.
- d_{ij} : The distance (or cost) between city i and city j .
- x_{ij} : A binary variable that indicates whether the route travels from city i to city j (1 if yes, 0 if no).

D. SIMULATED ANNEALING

Simulated Annealing is a probabilistic optimization algorithm inspired by the annealing process in metallurgy, which allows materials to reach a low-energy crystalline state through controlled heating and cooling. By making random moves and accepting worse solutions with a certain probability—reflected in the change in evaluation value ΔE —the algorithm effectively balances exploration and exploitation in complex search spaces. As the algorithm progresses, the likelihood of accepting inferior moves decreases, mimicking the cooling process and enabling the search of both promising and sub-optimal regions, thereby facilitating the discovery of high-quality solutions in complex optimization problems across various fields, including operations research, engineering, and artificial intelligence.

Algorithm:

In this context, ΔE represents the energy difference, and we assess the probability of making a move based on this value within the conditional check. The variable T denotes the temperature, which starts at a high value and gradually decreases through a cooling function. Initially, all moves are equally likely due to this high temperature.

```

current ← INITIAL - STATE
T ← some large value
for t = 1 to ∞ do
    if termination criteria then
        break
    end if
    next ← a randomly selected successor of current
    ΔE ← VALUE(next) - VALUE(current)
    if ΔE > 0 then
        current ← next
    else
        current ← next only with probability 1/(1 + e-ΔE/T)
    end if
    T ← cooling - function(T)
end for

```

E. PROBLEM STATEMENTS

- 1) Given a list of at least twenty important tourist locations in Rajasthan, and assuming the cost of travel between two locations is proportional to the distance between them, use the simulated annealing algorithm to plan a cost-effective tour. The goal is to find a tour that visits each location exactly once and returns to the starting point, minimizing the total travel cost.

Problem Definition:

- **Input:** A graph representing cities as nodes, with the distances between them represented as weighted edges.
- **Output:** A route that goes to each city only one time and comes back to the starting city, while keeping the total travel cost (or distance) as low as possible.

Start-up:

- 1) Identify 20 significant tourist locations in Rajasthan to serve as nodes in the graph.
- 2) Set the initial temperature T to a high value (e.g., $T = 1000$).
- 3) Define a minimum temperature T_{\min} to terminate the annealing process (e.g., $T_{\min} = 10$ or $T_{\min} = 102$).
- 4) Establish a cooling schedule by reducing the temperature by a fraction α after each iteration.
- 5) Generate an initial tour by creating a random permutation of the selected cities.
- 6) Calculate the cost (distance) of the initial tour using the distances between the identified cities.

Simulated Annealing

- 1) Repeat the following steps until the temperature reaches T_{\min} :
 - a) **Generate Neighbor:** Create a neighboring tour by perturbing the current tour. This may include:
 - Swapping two cities in the tour.
 - Reversing a sequence of cities in the tour.
 - Shifting a subset of cities within the tour.
 - b) **Calculate Cost:** Determine the total distance of the new tour.
 - c) **Acceptance Probability:** Compute the acceptance probability using the formula:

$$\text{Acceptance Probability} = e^{-\frac{\Delta f}{T}},$$

where Δf is the difference in cost between the new tour and the current tour (new cost - current cost), and T is the current temperature.

- d) **Accept or Reject Neighbor:** Accept the new tour based on the acceptance probability. Always accept the new tour if it has a lower cost than the current tour.
- e) **Update Tour:** If the new tour is accepted, update the current tour to the new tour and adjust the current cost accordingly.
- f) **Cooling:** Decrease the temperature following the established cooling schedule.

- 2) Track the best tour and its corresponding distance found throughout the iterations.

Solution:

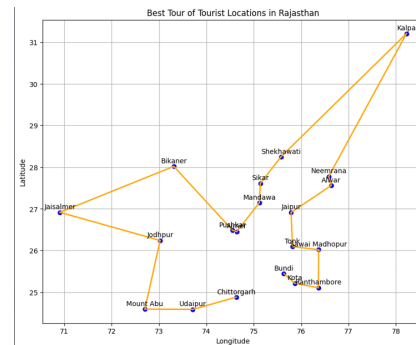


Figure 5: Best Tour Path.




Figure 6: Cost vs Iteration




Figure 7: Acceptance Probability.

- 2) You need to solve a jigsaw puzzle represented as a 512x512 pixel image. The image is divided into smaller pieces and the objective is to rearrange these pieces to reconstruct the original image.

Problem Definition:

- **Input:** The input consists of a file named `scrambled_lena.mat`, which contains the scrambled image data represented as a matrix of pixel values. The image dimensions are 512x512 pixels, and it can be divided into smaller square pieces, such as 64 pieces of 128x128 pixels each. The initial configuration is represented by a random arrangement of the jigsaw pieces, serving as the starting state for the simulated annealing algorithm.
- **Output:** The desired output is the reconstructed image, which should reflect the solved configuration of the original Lena image. This reconstructed image will be properly arranged and can be saved or displayed as a 512x512 pixel image, representing the final state after the simulated annealing process has been completed.

Simulated Annealing in Puzzle Solving:

1) **Initial State:** The initial state consists of a random arrangement of the tiles. Each tile represents a 128×128 pixel block of the image. This scrambled configuration serves as the starting point for the algorithm.

2) **Energy Calculation:** The **energy function** quantifies how "disordered" the current arrangement of tiles is. This is computed based on the differences in pixel values between adjacent tiles:

- **Left-Right Energy:** Measures pixel differences along the vertical edges of tiles.
- **Up-Down Energy:** Measures pixel differences along the horizontal edges of tiles.

The total energy is the sum of energy contributions from all tiles, reflecting how well the tiles fit together.

3) **Neighbor Generation: Neighboring Solutions:** The algorithm generates neighboring configurations by swapping two tiles. This represents a small perturbation in the current arrangement of pixels. Each swap results in a new arrangement of the tiles, leading to a different configuration of the overall image.

4) **Acceptance Probability:** The algorithm calculates the change in energy (ΔE) after a swap. If the new configuration has lower energy (i.e., a better fit), it is accepted. If the new configuration has higher energy (worse fit), it may still be accepted with a certain probability, which decreases as the algorithm progresses and temperature lowers. This allows for exploration of less optimal configurations to escape local minima.

5) **Cooling Schedule:** The temperature T starts high, allowing for more random swaps and exploration of the solution space. Over time, the temperature decreases, reducing the likelihood of accepting worse configurations, and focusing more on refining the best found solution.

6) **Final State:** The process continues until the temperature reaches a predefined minimum (T_{\min}), at which point the algorithm stops. The resulting arrangement of tiles represents the reconstructed image, ideally closely resembling the original.

Solution:




Figure 8: Acceptance Probability.




Figure 9: Energy Over Time.




Figure 10: Image Flattened



Figure 11: Solved Puzzle

F. Conclusion

In conclusion, this lab report explores the application of the simulated annealing algorithm to solve the problem of unscrambling an image represented by the `scrambled_lena.mat` file. By dividing the 512x512 pixel image into 64 smaller pieces of 128x128 pixels, we established a random initial configuration that served as the starting point for our algorithm. Through iterative adjustments and temperature-based heuristics, the simulated annealing process effectively optimized the arrangement of the jigsaw pieces, gradually leading to the correct configuration of the image. This project highlights the robustness of simulated annealing in solving complex optimization problems, demonstrating its ability to navigate through various configurations to achieve a desired outcome. The results not only illustrate the potential of this approach in image processing but also contribute to a deeper understanding

of the principles behind heuristic search algorithms in artificial intelligence.

VI. WEEK 5: GRAPHICAL MODELS AND BAYESIAN NETWORKS INFERENCE AND CLASSIFICATION IN R

A. Introduction

In this lab, we explore how to use graphical models and Bayesian inference to handle uncertainty in data. The focus is on educational data, where we will build Bayesian Networks using R to model the relationships between student grades in different courses. We also explore classification, specifically using the naive Bayes method, which is a simple but effective way to classify data based on probability. The dataset includes student grades and their qualification status for an internship. The goal is to train and evaluate the naive Bayes classifier, both by assuming the grades are independent and by considering the dependencies between them. This lab serves as an introduction to using probabilistic models for data analysis and classification.

B. Fundamentals

1) **Bayesian Inference and Graphical Models:** Graphical models use simple diagrams to show how different variables are related, which helps us understand uncertainty better. Bayesian inference is a method that lets us change our beliefs about unknown things when we get new information. Together, these ideas help us make sense of complex data and make better decisions.

Before starting practical work, it's important to understand the basics of graphical models and Bayesian inference. This includes learning about conditional probability, Bayes' theorem, and how graphical models show relationships between variables. These ideas help us make better decisions when dealing with uncertainty.

2) **Constructing Bayesian Networks in R:** Bayesian Networks are types of graphical models that show how different variables are related using a directed acyclic graph (DAG), which is a diagram that connects the variables without any cycles. In R, there are packages like 'bnlearn' that make it easy to build, visualize, and analyze these Bayesian Networks.

Learners will understand how to create the structure of Bayesian Networks, set up Conditional Probability Tables (CPTs), and carry out tasks like asking probabilistic questions and learning from data.

3) **Learning Dependencies from Data:** Bayesian Networks often rely on expert insights to determine how variables are related, but these relationships can also be learned from data. Methods like structure learning and parameter learning help uncover these connections and build Conditional Probability Tables (CPTs) from observed data. Participants will learn how to use data to identify the structure and parameters of Bayesian Networks, exploring algorithms such as constraint-based methods, score-based methods, and hybrid approaches to discover these relationships

4) **Naive Bayes Classification:** Naive Bayes is a straightforward probabilistic model that applies Bayes' theorem while assuming that the features are independent of one another. Although it is a basic method, it is commonly employed in applications such as document classification, sentiment analysis, and email filtering.

Students will comprehend the underlying principles of naive Bayes classification, including the calculation of class probabilities using Bayes' theorem and the assumption of feature independence.

Bayes Theorem:

$$P(C|X) = \frac{P(X|C) \cdot P(C)}{P(X)}$$

Naive Bayes Classification Formula:

$$P(C|X) \propto P(C) \cdot P(x_1|C) \cdot P(x_2|C) \cdot \dots \cdot P(x_n|C)$$

Where:

- $P(C|X)$ is the **posterior probability**: the probability of class C given the features X .
- $P(X|C)$ is the **likelihood**: the probability of the features X given the class C .
- $P(C)$ is the **prior probability** of the class C .
- $P(X)$ is the **marginal likelihood** or evidence: the total probability of the features X .

5) **Classifier Implementation and Evaluation:** Setting up a Naive Bayes classifier involves training the model on labeled data by calculating class priors (how often each outcome occurs) and conditional probabilities (how likely each feature is for each outcome). Once trained, the model can predict outcomes for new data based on these probabilities. To assess the model's performance, you use evaluation metrics like accuracy, precision, recall, and F1 score. Learners will gain practical experience using R to implement this classifier, covering data preparation, model training, prediction, and evaluation, ultimately applying these techniques to real-world educational data analysis tasks

C. Problem Statement

1) **Problem I:** The dataset `2020_bn_nb_data.txt` includes student grades from various courses. The objective is to model the relationships between these courses and learn the Conditional Probability Tables (CPTs). Furthermore, the aim is to predict a student's grade in PH100 based on their grades in other courses.

We will use the last column, which indicates internship qualification, to train a Naive Bayes classifier with 70% of the data. This classifier will predict if a student qualifies for an internship based on their grades, assuming that the courses are independent. We'll evaluate its accuracy on the remaining 30% of the data over 20 runs. Afterward, we will repeat the experiment, this time considering any potential relationships between the grades earned in different courses..

D. Implementation Method

1) Data Preprocessing:

- Load the dataset (*2020_bn_nb_data.txt*) into R.
- Analyze the data to gain insights into its structure and content.
- Clean the data as needed by addressing any missing values or outliers.

2) Bayesian Network Construction:

- Use the `bnlearn` package to construct a Bayesian Network.
- Establish the network structure using domain knowledge or apply structure learning algorithms to derive it from the data.
- Estimate Conditional Probability Tables (CPTs) for each node in the network.

3) Grade Prediction in PH100:

- Given a student's grades in other courses, use the Bayesian Network to predict the grade in PH100.
- Use the known grades to query the network and determine the most probable grade in PH100.

4) Naive Bayes Classifier:

- Split the dataset into training and testing sets (70% training, 30% testing).
- Train a naive Bayes classifier using the training data, assuming independence between course grades.
- Evaluate the classifier's performance on the testing data using accuracy and other relevant metrics.

5) Assessing Classifier Accuracy:

- Repeat the training and testing of the Naive Bayes classifier 20 times using randomly selected data.
- Document the accuracy of the classifier during each iteration.

6) Considering Dependencies:

- Adjust the Naive Bayes classifier to account for possible relationships between course grades.
- Redo the training and testing process, evaluating the accuracy of the classifier based on this revised approach.




Figure 12: Dependencies

F. Conclusion

In this lab, we learned about graphical models, Bayesian inference, and classification within the realm of educational data analysis. By building Bayesian Networks, we were able

```
> library(bnlearn)
> library(caret)
Loading required package: ggplot2
Loading required package: lattice
> library(e1071)
> grades <- c("AA", "AB", "BB", "BC", "CC", "CD", "DD", "FF")
> course.grades <- read.table("C:/Users/Abhi Patel/Downloads/2020_bn_nb_data.txt", header=TRUE)
> set.seed(100)
> tIndex <- createDataPartition(course.grades$QP, p=0.7, list=FALSE)
> train <- course.grades[tIndex, ]
> test <- course.grades[-tIndex, ]
> nbc <- naiveBayes(QP ~ EC100 + EC160 + IT101 + IT161 + MA101 + PH100 + PH160 + SS101, data=train)
> printAll <- function(model) {
+   trainPred <- predict(model, newdata = train, type = "class")
+   trainTable <- table(trainQP, trainPred)
+   trainAcc <- sum(diag(trainTable)) / sum(trainTable)
+   testPred <- predict(model, newdata = test, type = "class")
+   testTable <- table(testQP, testPred)
+   testAcc <- sum(diag(testTable)) / sum(testTable)
+   message("Accuracy")
+   print(round(cbind("Training Accuracy" = trainAcc, "Test Accuracy" = testAcc), 4))
+ }
> printAll(nbc)
Accuracy
Training Accuracy Test Accuracy
[1,]
0.9939 0.9882
> |
```

Figure 13: Final Solution

to model how student grades in different courses relate to one another and create Conditional Probability Tables (CPTs) from the data. We also used a Naive Bayes classifier to predict whether students qualify for a internships based on their grades, initially assuming that course grades are independent. Through our experiments, we found that while the classifier performed well without considering dependencies, taking these relationships into account improved its accuracy and better reflected real-world situations. Overall, this lab emphasized the importance of using graphical models and understanding dependencies in order to make more accurate predictions and informed decisions in educational contexts.

VII. CHALLENGE PROBLEM : GENERATING MELODY USING SIMULATED ANNEALING

A. Problem Statement

Create a computer program to generate a melody in the style of North Indian Classical Raag Bhairav music. Use advanced algorithms (simulated annealing or genetic algorithms) to ensure the generated melody follows the traditional rules and patterns of Raag Bhairav, including key phrases typically used in this style of music

B. Key Components

1) Sargam and Note Mapping:

- Sargam Notes: Sa, Re, Ga, Ma, Pa, Dha, Ni

Sargam	MIDI Note
Sa	60
Re	62
Ga	64
Ma	65
Pa	67
Dha	69
Ni	71

- MIDI Note Mapping:

2) Target Phrase:

Target Phrase: Sa Re Ga Pa Ma Pa Dha Ni Sa

3) Simulated Annealing:

- Mutation: Randomly changes a note in the melody.
- Acceptance Probability: Determines the likelihood of accepting a worse solution based on temperature.

4) *MIDI File Generation:*

- Uses a MIDI library (e.g., 'mido') to create a MIDI file from the generated melody.

C. *Simulated Annealing Algorithm in Melody Generation*

1) *Key Parameters:*

- P_i : Initial phrase
- P_t : Target phrase
- T_i : Initial temperature
- T_{min} : Minimum temperature
- α : Cooling rate
- $S(P)$: Score of phrase P
- ΔS : Change in score

2) *Iterative Process:*

- 1) Temperature Update:

$$T = \max(T_i \cdot \alpha^i, T_{min})$$

- 2) Mutation:

$$P_{new} = M(P_i)$$

- 3) Score Calculation:

$$\Delta S = S(P_{new}) - S(P_i)$$

- 4) Acceptance:

$$\text{Accept } P_{new} \text{ with probability } \min\left(1, \exp\left(\frac{\Delta S}{T}\right)\right)$$

3) *Scoring Function:*

$$S(P) = \sum_{i=1}^n 2 \cdot \delta(P_i, P_t[i]) + \delta(P_i \in P_t, 1)$$

where $\delta(x, y)$ is the Kronecker delta:

$$\delta(x, y) = \begin{cases} 1, & \text{if } x = y \\ 0, & \text{otherwise} \end{cases}$$

D. *Conclusion*

The proposed Simulated Annealing-based melody generator effectively produces high-quality Raag Bhairav melodies, adhering to traditional grammatical rules and structures. By leveraging advanced mutation techniques and a well-defined scoring function, the algorithm efficiently explores the search space. This work lays the groundwork for AI-generated Indian classical music research, with future directions including: exploring alternative optimization algorithms (genetic, particle swarm), expanding the dataset for improved generalization, applying the approach to other North Indian ragas, and integrating human expertise for enhanced authenticity. This innovative solution demonstrates the potential of AI in preserving and innovating within traditional musical heritage, paving the way for further experimentation and development in music generation.

VIII. *CONCLUSION*

Our experience in the Artificial Intelligence Lab has provided us with critical insights and hands-on expertise in the realm of efficient AI techniques. By exploring diverse methods such as state-space search, heuristic optimization, probabilistic reasoning, and machine learning, we have developed a strong foundation in problem-solving strategies within AI. Through practical experiments and implementation of various algorithms, we have gained a comprehensive understanding of how AI can be applied to solve complex, real-world problems. This experience has equipped us with the knowledge and skills to adapt and thrive in the ever-evolving field of artificial intelligence, preparing us for future challenges and advancements in the discipline.

IX. *ACKNOWLEDGEMENT*

We would like to extend our gratitude to Dr. Pratik Shah for his invaluable guidance, motivation and support throughout the development of this lab manual. His expertise and mentorship have been crucial in shaping our understanding of the subject. We greatly appreciate his time and contributions to our academic journey.

X. *REFERENCES*

- Russell, S., & Norvig, P. (2020). Artificial intelligence: A modern approach (4th ed.). Pearson
- Khemani, D. (2014). A first course in artificial intelligence. Tata McGraw Hill
- Masrat, A., Gawde, H., Makki, M.A., & Parekh, U. (2021). Document Plagiarism Detection Tool using Edit Distance Text Similarity Measure. International Research Journal of Engineering and Technology (IRJET), 08(09), 1395-1399.
- Gardahadi, G. (2019). Plagiarism Detection Using Levenshtein Distance With Dynamic Programming.

XI. *GITHUB REPOSITORY*

For additional resources and detailed codes, please refer to our GitHub repository [link](#). This repository contains the source code, implementation details, and contributions from our team.